Programming Assignment 2

A Simple Gnutella-style P2P File Sharing System

Submission:
- *Due by 11:59pm of 02/26/2017.*
- *Late penalty: 20% penalty for each day late.*
- *You may work individually or in groups of two for this assignment. For the groups with 2 members, only one submission listing both members is needed.*
- *Please upload your assignment on the Blackboard with the following name: Section_ LastName_FirstName_PA2.*
- *Please do NOT email your assignment to the instructor and TA!*

1   The problem

In programming assignment 1, you implemented a Napster style file-sharing system where a central indexing server plays an important role. In this project, you are going to remove the central component and implement a pure distributed file-sharing system. An example of such a system is the well-known **Gnutella network**.

You can be creative with this project. *You are free to use any programming languages (e.g., C/ C++ or Java) and any abstractions such as sockets, RPCs, RMIs, threads, events, etc. that might be needed. Also, you are free to use any machines such as your laptops or PCs.*

If you are not familiar with Gnutella, the following links provide some background about the technology. Pay more attention to its architecture and design goals rather than its protocol details, since we are not strictly implementing a full-featured Gnutella client, but a small subset of it.

Gnutella Protocol at http://rfc-gnutella.sourceforge.net/developer/stable/index.html

In this programming assignment, you need to design **a Gnutella-style peer-to-peer (P2P) system**. Like in Programming Assignment 1, each peer should be both a server and a client. As a client, it provides interfaces through which users can issue queries and view search results. As a server, it accepts queries from other peers, checks for matches against its local data set, and responds with corresponding results. In addition, **since there's no central indexing server, search is done in a distributed manner.** Each peer maintains a list of peers as its neighbor. Whenever a query request comes in, the peer will broadcast the query to all its neighbors in addition to searching its local storage (and responds if necessary).

Below is a list of what you need to implement:

1. First of all, we are not implementing the dynamic initialization of the network as in Gnutella. To keep things simple, assume that the structure of the P2P network is **static**. This means you don't need to implement group membership messages (PING/PONG in Gnutella). Your network will be initialized statically using a config file that is read by each peer at startup time. You are free to use any format for your config file. The only requirement is that it provides a list of all neighbors for that peer.

2. Having initialized the P2P network, a peer searches for files by issuing a query. The query is sent to all neighbors. Each neighbor looks up the specified file using a local index and responds with a *queryhit* message in the event of a hit. The *queryhit* message is propagated back to the original sender by following the reverse path of the query (the following paragraph describes how this can be done). Regardless of a hit or a miss, the peer also forwards the query to all of its neighbors.

   To prevent query messages from being forwarded infinitely many times, each query message carries a *time-to-live (TTL)* value that is decremented at each hop. In addition, each query has a globally unique message ID (defined for our purposes as a [peer ID, sequence number]). Each peer also keeps track of the message IDs for messages it has seen, in addition with the upstream peers where the messages are sent from. You can use an associative array that stores [message ID, upstream peer ID] pairs, where the upstream peer ID could be a peer's IP address (and port number if necessary). Use your own heuristics to decide the size of this associative array your peer needs to maintain, and flush out old entries at appropriate times (in essence, you don't want this buffer to grow indefinitely). The purpose of maintaining this data structure at each peer is two-fold, one is to prevent a peer from forwarding a message it already saw (and forwarded), and the second is to provide the reverse path for the *queryhit* message to propagate back to the original sender of the query. Note that the *queryhit* message MUST carry the same message ID as the corresponding query in order to be propagated back correctly.

   Ideally *query* and *queryhit* messages can be propagated by using TCP socket connections. But you can also use RPCs or RMIs if you so wish.

   Assuming the above description, the message formats are defined as follows:

   - query (message ID, TTL, file name)
   - queryhit(message ID, TTL, file name, peer IP, port number)

3. Routing of messages is by broadcast/back-propagation manner. Each peer maintains a list of its neighboring peers (picked statically by you). A query message from S is broadcasted to all of S's neighbors and relayed by each receiver until its TTL value decreased to 0. A *queryhit* message is sent back to the original sender following the reverse path. The message ID is used for this purposes as described previously.

4. *Obtain* is achieved by sending a direct download request to a peer that sends a *queryhit* message back, this is pretty much the same as done in project one.

   - obtain (file name)

Other requirements:

- Use threads so your peer can serve multiple requests concurrently.
- No GUIs are required.

2   Evaluation and Measurement

Deploy at least 10 peers. They can be setup on the same machine (different directories) or different machines. Each peer has in its shared directory at least 10 text files of varying sizes (for example 1k, 2k, ..., 10k). Make sure some files are replicated at more than one peer sites (so your query will give you multiple results to select).

Do the following experiment in at least the following two topologies (initialize the topology by assigning neighbors for each peer):

- **A star topology.**
- **A linear topology**.

Do a simple experiment to evaluate the behavior of your system. Compute the average response time per client query request by measuring the average response time seen by a client, since there may be multiple results for each query, measure the average among them. And repeat this measurement for 200 times and get the average. Use your own judgment/technique to decide when the last query result should come back. For example, define a cutoff time, waiting until that time and compute the result.

Do the same calculation by changing system load, more specifically, do the same experiment where there's only 1 client issuing queries, then 2 clients, 3 clients, and so on. Draw a plot after collecting all the data and justify your conclusion. Also compare the result to the first programming assignment 1 and justify your conclusion.

You should compare the system performance between a star topology and linear topology under the same system size. List their advantages, disadvantages and applicability.

You may calculate the system load (in terms of generated traffics, bytes/sec). And make comments.

3   What you will submit

When you have finished implementing the complete assignment as described above, you should submit your solution on blackboard.

Each program must work correctly and be **detailed in-line documented**. You should hand in:

1. **Output file**: A copy of the output generated by running your program. When it downloads a file, have your program print a message "display file 'foo'" (don't print the actual file contents if they are large). When a peer issues a query (lookup) to the indexing server, having your program print the returned results in a nicely formatted manner.
2. **Design Doc**: A separate (**typed**) design document (named *design.pdf* or *design.txt*) of approximately 2-4 pages describing the overall program design, , and design tradeoffs considered and made. Also describe possible improvements and extensions to your program (and sketch how they might be made).
3. **Source code and program list**: All of the source code and a program listing containing in-line documentation.

4. **Manual:** A detailed manual describing how the program works. The manual should be able to instruct users other than the developer to run the program step by step. The manual should contain at least a test case which will generate the output matching the content of the output file you provided in 1.

5. **Verification**: A separate description (named *test.pdf* or *test.txt*) of the tests you ran on your program to convince yourself that it is indeed correct. Also describe any cases for which your program is known not to work correctly.

6. **Performance results.**

7. **Problem report:** if your code does not work or does not work correctly, you should report this.

Please put all of the above into one .zip or .tar file, and upload it on blackboard. The name of .zip or .tar should follow this format:

Section_LastName_FirstName_PA2.zip.

Please do **NOT** email your files to the professor and TA!!

4   Grading policy for all programming assignments

- Program Listing
    - works correctly ------------- 50%
    - in-line documentation -------- 15%
- Design Document
    - quality of design ------------ 15%
    - understandability of doc ------- 10%
- Thoroughness of test cases ---------- 10%

Grades for late programs will be lowered 20 points per day late.